

Session 3

Memory Structure and Pointer Manipulation

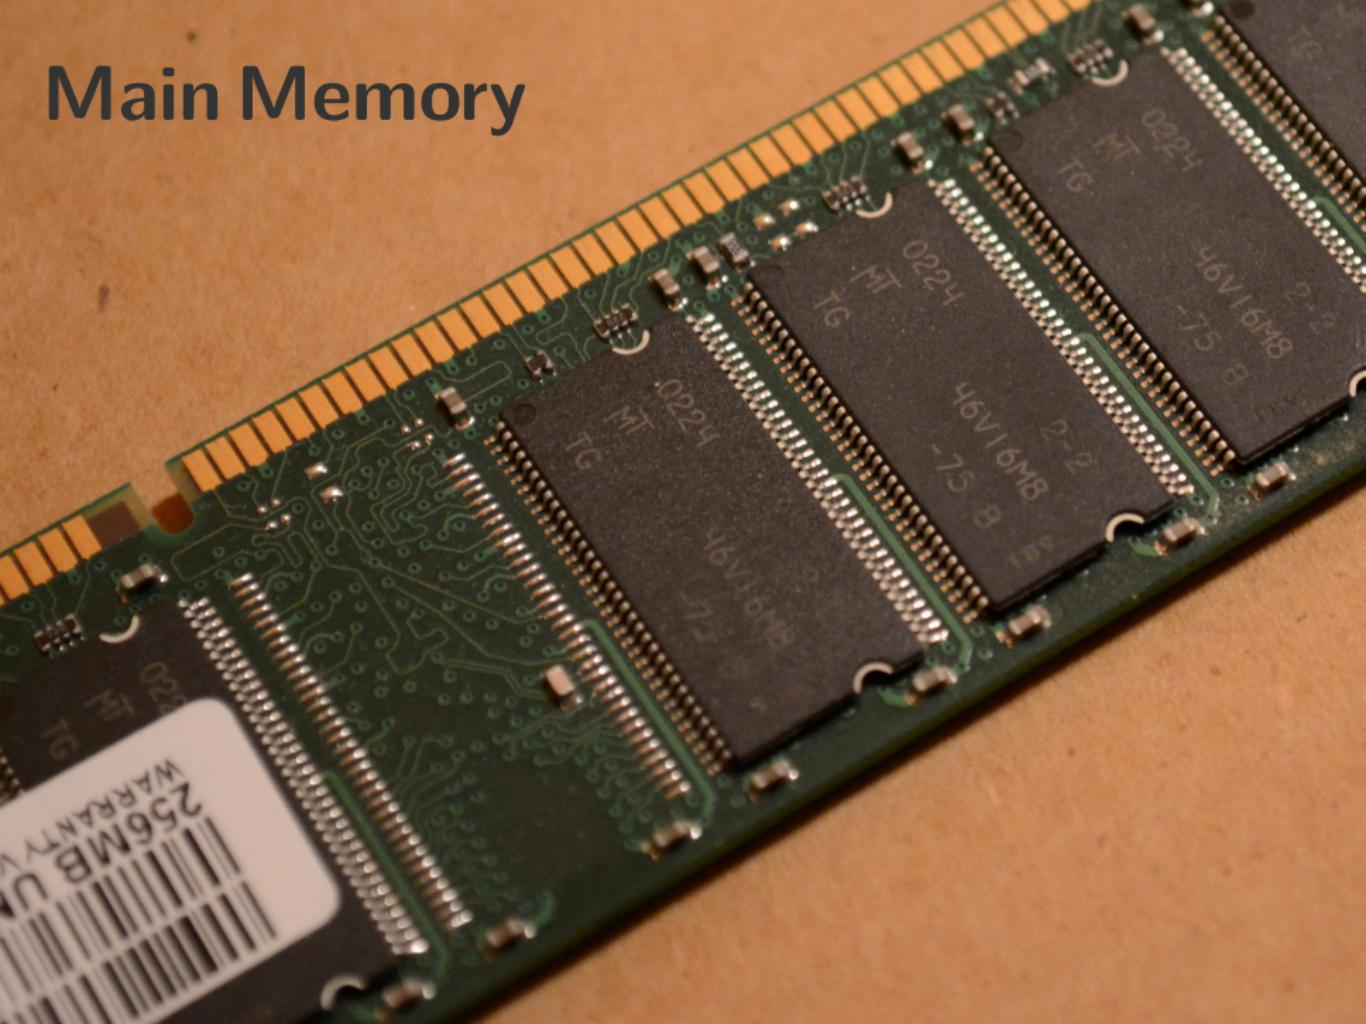


This work is licensed under a Creative Commons Attribution – NonCommercial – NoDerivatives 4.0 International License.

Objectives

- Processes and organisation of the **main memory**
 - Variable, addresses and pointers in the stack
 - Execution environment, procedure and function calls
 - Heap and creation of dynamic memory
- **Pointer arithmetic** and static/dynamic arrays
 - Pointer, variable address and dereferencing operator
 - Allocating memory with `malloc` and `free` functions

Main Memory



Main Memory (1)

- A running program is called a **process**
 - A program is a sequence of bit stored on the hard drive
 - A process is in the main memory and running
- Memory space used by a process split in **three parts**
 - The text of the program with all the instructions
 - The stack with all the local variables
 - The heap with all the dynamically allocated memory
- Memory is just a huge **bytes array**

Each byte has a unique address in memory

Main Memory (2)

- A **variable** is characterised by four elements
 - A memory address
 - The number of bytes used in memory
 - The value of the variable stored in memory
 - A symbolic name representing the variable

```
1 int a = 17;
```

2004	
a:	2000

Uninitialised Variable

- A variable just declared is said to be **uninitialised**
Its value is undefined and depends on the external factors
- The main memory is **not cleared** upon process creation

```
1 #include <stdio.h>
2
3 int main()
4 {
5     int a;
6     printf("%d\n", a);
7     return 0;
8 }
```

277119030

Addressing

- Amount of addressable memory depends on several elements
 - Number of bits used to store an address
 - Addressing granularity (one for each bit, nibble, byte, etc.)
- Some common examples
 - 32 bits address for each byte: $2^{32} = 4 \text{ GiB}$
Windows case that could not exceed 4 GiB of RAM
 - 64 bits address for each byte: $2^{64} = 16 \text{ EiB}$
We have time before seeing such RAM memory capacities

Endianness (1)

- **Endianness** defines the order in which bytes are organised

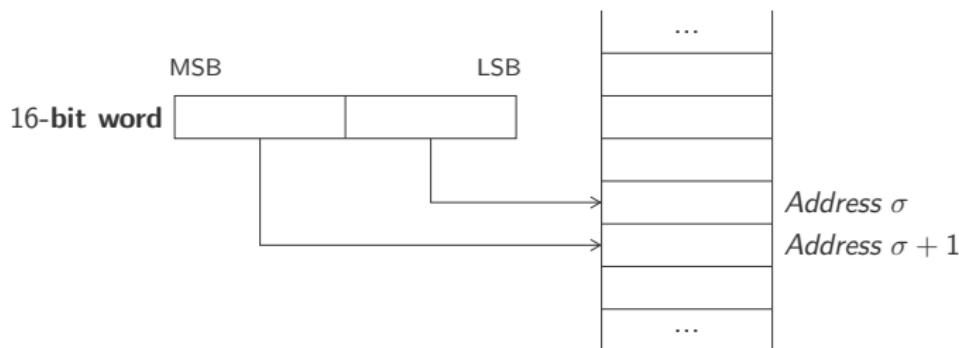
For a given word that is stored in memory



Endianness (2)

- Bytes are stored starting from LSB with **little endian**

The bytes are stored “in reverse order”





Environment Stack

Execution Environment

- Running procedure/function has its own **environment**

Contains memory space allocated for local variables

- **Sizes** known at compile time

Compiler can compute this size for each procedure/function

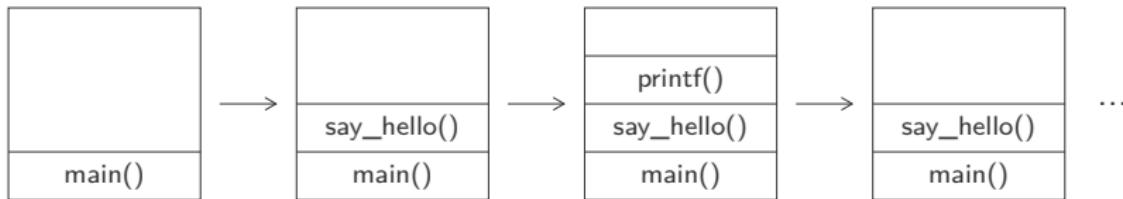
- Retain memory address where to put the **return value**

Only used for functions by the return statement

Environment Stack

- Saving current environment at each procedure/function call
 - To not loose local variables from calling procedure/function
 - Restore environment after running called procedure/function
- Using an environment **stack** in memory

Two operations on the stack: environment push and pop

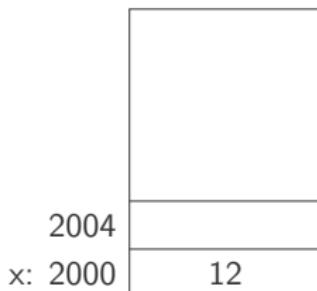


Procedure Call (1)

- The **entry point** of the program is the `main` function

A local variable is required for this procedure

```
1 int main()
2 {
3     int x = 12;
4     print_sum(x, -2);
5
6     return 0;
7 }
```



Procedure Call (2)

- Initialising parameters of print_sum during the call

Value of effective parameters copied in formal parameters

```
1 void print_sum(int a, int b)
2 {
3     int sum = a + b;
4     printf("%d + %d = %d\n", a, b, sum);
5 }
```

2016	
sum:	2012
b:	2008 -2
a:	2004 12
x:	2000 12

Procedure Call (3)

- Execution of the body of the print_sum procedure

The only accessible memory is the local environment of the call

```
1 void print_sum(int a, int b)
2 {
3     int sum = a + b;
4     printf("%d + %d = %d\n", a, b, sum);
5 }
```

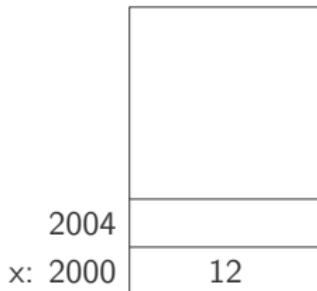
2016	
sum:	2012
b:	2008
a:	2004
x:	2000

Procedure Call (4)

- After execution of `print_sum`, **return to main function**

The environment of the `print_sum` call is deleted

```
1 int main()
2 {
3     int x = 12;
4     print_sum(x, -2);
5
6     return 0;
7 }
```



Stack Overflow

- Limitation of the memory area reserved for the stack

The number of successive calls of procedures/functions is limited

- The precise limitation depends on the machine

May depend on some physical constraints (register, memory, etc.)

```
1 int main()
2 {
3     main();
4     return 0;
5 }
```

```
[1]      70751 segmentation fault  ./a.out
```

Pointer



東洋緑化(株)

3M

Pointer (1)

- A **pointer** is a variable containing an address
 - Unary operator & to get the address of a variable
 - Declaring a pointer with a * before the variable name
- A pointer is a **variable pointing** to another one

It stores a reference to a memory cell

```
1 char c = 'M';
2 char *p = &c;
```

2009	
p:	2001
c:	2000

'M'

Dereferencing Operator (1)

- Dereference **unary operator *** (or indirection operator)

Dereferencing a pointer means accessing to the pointed variable

- Dereferencing pointer type `xxx *` gives value with type `xxx`

Otherwise, it produces a compile error

```
1 short a = 11;
2 short *p = &a;
3 short b = *p;
```

2012	
b:	11
p:	2002
a:	2000

Pointer (2)

- It is possible to store the **address of a pointer** in a pointer

Which gives two levels of indirection

```
1 short a = 11;
2 short *p = &a;
3 short b = *p;
4
5 short **pp = &p;
```

2020	
pp:	2012
b:	2010
p:	2002
a:	2000

Pointer Type

- The *** symbol** indicates a pointer in the type of the variable

Mandatory to indicate what type the pointer refers to

- The `value` variable is a **pointer to an int**

```
int *value;
```

- The `value` variable is a **pointer to an int**

```
int **value;
```

Dereferencing Operator (2)

- Modifying pointed variable value with dereferencing operator

Just use it to the left of the assignment operator

```
1 short a = 11;  
2 short *p = &a;  
3 *p = 99;
```

2010	
p:	2002
a:	2000

Endianness (3)

- Example with a `short int` variable stored at address σ

Accessing separately the two bytes with a cast

```
1 #include <stdio.h>
2
3 int main()
4 {
5     short s = 277; // 256 + 21
6     char *c = (char*) &s;
7
8     printf("%d\n", *c);
9     printf("%d\n", *(c + 1));
10
11    return 0;
12 }
```

```
21
1
```

Function Call (1)

- The **entry point** of the program is the `main` function

Two local variables is required for this function

```
1 int main()
2 {
3     int x = 12;
4     int sum = get_sum(x, -2);
5
6     return 0;
7 }
```

2008	
sum:	2004
x:	2000 12

Function Call (2)

- Initialising parameters of get_sum during the call

Storing return value address in a special variable \$

```
1 int get_sum(int a, int b)
2 {
3     int sum = a + b;
4     return sum;
5 }
```

2028	
sum:	2024
b:	2020 -2
a:	2016 12
\$:	2008 2004
sum:	2004
x:	2000 12

Function Call (3)

- Execution of the first instruction of the function get_sum

The local variable sum is initialised with the result of a + b

```
1 int get_sum(int a, int b)
2 {
3     int sum = a + b;
4     return sum;
5 }
```

2028	
sum:	2024
b:	2020
a:	2016
\$:	2008
sum:	2004
x:	2000

Function Call (4)

- After execution of `get_sum`, return to `main` function

The return statement copies the return value (`$ = sum`)*

```
1 int main()
2 {
3     int x = 12;
4     int sum = get_sum(x, -2);
5
6     return 0;
7 }
```

2008	
sum:	2004
x:	2000

Passing Parameter by Reference (1)

- Parameters **passed by value** when calling procedure/function

Value of effective parameters copied in formal parameters

```
1 void swap(int a, int b)
2 {
3     int tmp = a;
4     a = b;
5     b = tmp;
6 }
7
8 int main()
9 {
10    int x = 2, y = 15;
11    swap(x, y);
12    printf("x = %d et y = %d\n", x, y);
13
14    return 0;
15 }
```

Passing Parameter by Reference (2)

- Passing by reference if parameters are pointers

Called procedure/function accesses values through addresses

```
1 void swap(int *a, int *b)
2 {
3     int tmp = *a;
4     *a = *b;
5     *b = tmp;
6 }
7
8 int main()
9 {
10    int x = 2, y = 15;
11    swap(&x, &y);
12    printf("x = %d et y = %d\n", x, y);
13
14    return 0;
15 }
```

Passing Parameter by Reference (3)

- Execution of the first instruction of the swap procedure

Local variable tmp initialised with value pointed by a

```
1 void swap(int *a, int *b)
2 {
3     int tmp = *a;
4     *a = *b;
5     *b = tmp;
6 }
```

2020	
tmp:	2016
b:	2012
a:	2008
y:	2004
x:	2000

Returning Pointer (1)

- Function that **returns a pointer** to a new variable

Declaring and initialising a variable with a specified value

```
1 int* new_int(int value)
2 {
3     int a = value;
4     return &a;
5 }
6
7 int main()
8 {
9     int *result = new_int(4);
10    printf("%d\n", *result);
11    return 0;
12 }
```

Returning Pointer (2)

- Execution of the first instruction of the swap procedure

Local variable tmp initialised with the value pointed by a

```
1 int* new_int(int value)
2 {
3     int a = value;
4     return &a;
5 }
```

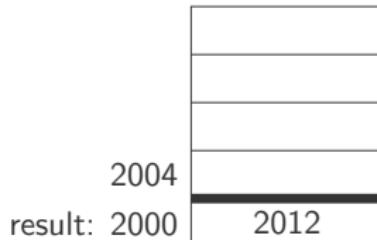
2020	
a: 2016	4
value: 2012	4
\$: 2004	2000
result: 2000	

Returning Pointer (3)

- Return from the call and **deletion of associated environment**

The pointer result of main does not point anything

```
1 int main()
2 {
3     int *result = new_int(4);
4     printf("%d\n", *result);
5     return 0;
6 }
```



Array



Array

- An **array** can store a sequence of elements
 - All the elements have the same type (int, double, char, etc.)
 - A size N defining the number of elements
- Accessing the **cell of the array** with index i with $\text{arr}[i]$
Indices of array cells goes from 0 to $N - 1$
- **Example** of an array arr with size 6



Array Operation

- Declaring an array with `int arr[N];`

Specifying the type of elements and the size of the array

- Modification of an element with the `assignment operator`

For example, `arr[i] = 3;` stores the value 3 in the index i cell

```
1 int N = 3;                      // Size of the array
2 int arr[N];                     // Declaration of an array of integers
3
4 int i;
5 for (i = 0; i < N; i++)
6     arr[i] = i + 1;             // Store values in the array
```

Array Initialisation

- Array cells values **not initialised** at declaration

Manual initialisation or with an initialiser

- **Direct initialisation** of an array with initial values

Only possible when declaring the array, not after

```
1 int a[3] = {3, 2, 1};           // Declaration and initialisation
2 a = {1, 2, 3, 4, 5}           // => FORBIDDEN
3
4 int b[10] = {1, 2};            // Specified size is an upper bound
5 int c[] = {1, 2, 3, 4, 5};     // Size can be automatically inferred
```

Displaying Array Elements

- Display elements of an array arr with a loop

Aesthetic display in brackets, separating elements with commas

- Computing the size of the array with the sizeof operator

sizeof(type) for type, sizeof var for variable, in bytes

```
1 int arr[] = {5, 4, 3, 2, 1};
2 int N = sizeof arr / sizeof(int);
3
4 if (N == 0)                                // Base case, empty array
5     printf("[]");
6 else {
7     printf("[%d", arr[0]);                  // Display 1st element
8
9     int i;
10    for (i = 1; i < N; i++)
11        printf(", %d", arr[i]);           // If several elements
12    printf("]\n");
13 }
```

Array in Memory

- An array is a **pointer** to memory area of contiguous cells
 - Stored in the procedure/function environment in the stack
 - Constant pointer whose value cannot change

```
1 int a[3] = {3, 2, 1};  
2  
3 int *p = a;  
4 printf("%d\n", p[0]);  
5 a = p; // => FORBIDDEN
```

2020	
p:	2012
	2000
2008	1
2004	2
a:	2000
	3

Array Parameter (1)

- Pass an **array parameter** to a procedure/function

Passing by reference, that is, as for a `int`*

```
1 void print_arr(int arr[])
2 {
3     int N = sizeof arr / sizeof(int);
4
5     // [...]
6 }
```

```
program.c:5:20: warning: sizeof on array function parameter will
      return size of 'int *' instead of 'int []'
      [-Wsizeof-array-argument]
          int N = sizeof arr / sizeof(int);

program.c:3:20: note: declared here
void print_arr(int arr[])
```

Array Parameter (2)

- The `sizeof` operator cannot access the **stack of the caller**
`sizeof arr` returns the size of a `int` (8 on a 64 bits machine)*
- The `print_arr` function only displays the **two first** elements
The value of `sizeof arr / sizeof(int)` is $8/4 = 2$

```
1 int main()
2 {
3     int a[3] = {3, 2, 1};
4     print_arr(a);
5
6     return 0;
7 }
```

```
[3, 2]
```

Array Parameter (3)

- Required to pass the **size of the array** in the parameters

*Two similar notations: int arr[3] or int *arr*

```
1 void print_arr(int arr[], int N)
2 {
3     printf("%p\n", arr);
4     // [...]
5 }
6
7 int main()
8 {
9     int a[3] = {3, 2, 1};
10    printf("%p\n", a);
11    print_arr(&a, 3);
12
13    return 0;
14 }
```

```
0x7fff5c96d84c
0x7fff5c96d84c
[3, 2, 1]
```

Array Parameter (4)

- Passing a pointer to a **fixed size array** as a parameter

Access to array cells values by dereferencing the received pointer

```
1 void print_arr(int (*arr)[3])
2 {
3     int N = sizeof(*arr) / sizeof(int);    // Access to referenced array
4     if (N == 0)
5         printf("[]\n");
6     else {
7         printf("[%d", (*arr)[0]);           // Access to referenced array
8
9         int i;
10        for (i = 1; i < N; i++)
11            printf(", %d", (*arr)[i]);      // Access to referenced array
12        printf("]\n");
13    }
14 }
15
16 int main()
17 {
18     int a[3] = {3, 2, 1};
19     print_arr(&a);
20
21     return 0;
22 }
```

Careful with Types!

- Do not confuse pointers array and array pointeur

- Pointers array: int *p[2];
- Array pointer: int (*q)[] ;

```
1 int main()
2 {
3     int i = 12;
4
5     int *p[2];
6     p[0] = &i;
7     printf("%d\n", *p[0]);
8
9     int data[] = {99, 98, 97};
10
11    int (*q) [];
12    q = &data;
13    printf("%d\n", (*q)[0]);
14 }
```

Returning Array

- Impossible for a function to **return an array**

Created on the stack and deleted when returning from function

- A function can return a **subarray**

Returning the pointer to the start cell of the subarray

```
1 int* sub_arr(int *arr, int start)
2 {
3     return &arr[start];
4 }
5
6 int main()
7 {
8     int a[] = {1, 2, 3, 4, 5};
9     int *b = sub_arr(a, 2);
10    print_arr(b, 2);           // Displays "[3, 4]"
11
12    return 0;
13 }
```



MACHINE à CALCULER de Léon BOLLÉE AU MANS. B^{TÉ} S.G.D.G. 1889
EXPOSITION UNIVERSELLE de PARIS 1889 MÉDAILLE D'OR

Pointer Arithmetic

Array and Pointer (1)

- Notions of **pointer** and **array** are very close

Both are used to identify memory areas

```
1 int main()
2 {
3     int a = 1;
4     int b = 2;
5     int c = 3;
6
7     int *p = &c;
8     int i;
9     for (i = 0; i < 3; i++)
10        printf("%d\n", p[i]);
11
12    return 0;
13 }
```

```
3
2
1
```

Array and Pointer (2)

- An array is a **block of memory cells** that are contiguous

An array variable is a pointer to the beginning of the block

```
1 int arr[] = {1, 2};  
2  
3 printf("%p\n", arr);  
4  
5 printf("%d\n", arr[0]);  
6 printf("%d\n", *arr);
```

```
0x7fff58e7d840  
1  
1
```

2008	
2004	2
arr:	1

Notation Equivalence (1)

- Equivalent notations between array and pointers

Using brackets and indices, or direct computation

- Example given the `int arr[];` declaration

`arr[i]`

`*(&arr + i)`

`&arr[i]`

`arr + i`

Pointer Arithmetic (1)

- Pointer arithmetic proposes several operations

The + operator is used to move in the memory

- Two main ways to move in an array
 - Navigation with the [] notation and an index
 - Navigation with the pointer and pointer arithmetic

```
1 for (i = 0; i < N; i++)  
2     printf("%d\n", tab[i]);
```

```
1 for (i = 0; i < N; i++)  
2     printf("%d\n", *(tab + i));
```

+ Operator

- The + operator moves the pointer forward of **one memory cell**
 - Bytes increment to the address depends on the pointer type
 - No control on the value of the pointer

```
1 short s = 666;
2 char *p = (char*) &s;
3
4 printf("%d\n", *p);           // Display "-102"
5 printf("%d\n", *(p + 1));    // Display "2"
```

0000001010011010
 $\overbrace{\quad\quad\quad}^{*(p + 1)}$ $\overbrace{\quad\quad\quad}^{*p}$

Moving in Memory (1)

- The + operator goes to the next “cell”

Of the memory area indicated by the pointer

```
1 int arr[] = {1, 2, 3, 4};  
2  
3 printf("%d\n", *arr);           // arr is 2000  
4 printf("%d\n", *(arr + 3));    // arr + 3 is 2012 (2000 + 3 * 4)
```

2016	
2012	4
2008	3
2004	2
arr: 2000	1

Moving in Memory (2)

- Possible to **modify values** in memory where we want

Just use the dereferencing operator on a pointer

```
1 int arr[4];
2
3 int i;
4 for (i = 0; i < 4; i++)
5     *(arr + i) = 4 - i;
```

2016	
2012	1
2008	2
2004	3
arr:	2000
	4

Pointer Arithmetic (2)

- Three categories of **operations possible** on pointers
 - Addition and subtraction on a pointer with + and -
 - Increment (decrement) a pointer with ++ (--)
 - Comparing pointers with ==, <, >, <= and >=
- Browse an array with **start and end pointers**

End pointer can be computed given the size of the array

```
1 int tab[] = {1, 2, 3, 4, 5};  
2 int *start = &tab[0];  
3 int *end = start + 4;  
4  
5 while (start <= end)  
6     printf("%d\n", *start++);
```

Segmentation Fault

- You must pay attention to **illegal access** to memory

Each process has its own memory address space

- Memory accesses** are controlled by the operating system

Thanks to a hardware component, namely the MMU

```
1 int a = 999;
2 int b = 0;
3
4 printf("%d\n", *(&b + 1));    // Displays "999"
5 printf("%d\n", *(&b + 500)); // 9679 segmentation fault ./a.out
```

Operator Priority

- You have to pay attention to the **operator priorities**

*In particular the priority of the dereferencing operator **

- Example** given the declaration `int *p;`

$(*p)++$

$*(p++)$

$*p + 2$

$*(p + 2)$

NULL Pointer

- Only **integer literal value** that can be used for a pointer is 0
 - Indicates the absence of a pointed memory area
 - Represented by the **NULL** constant
- Possible to **test the NULLity** of a pointer

Using the == and != operators

```
1 int *p = NULL;  
2  
3 if (p != NULL)  
4     printf("Valeur de p: %d\n", *p);
```

Dynamic Size Array (1)

- Using NULL to indicate the end of an array

Array of pointers in which the last one is followed by a NULL

```
1 int main()
2 {
3     int a = 6;
4     int b = -2;
5
6     int *arr[10] = {&a, &b};
7     arr[2] = NULL;
8
9     return 0;
10 }
```

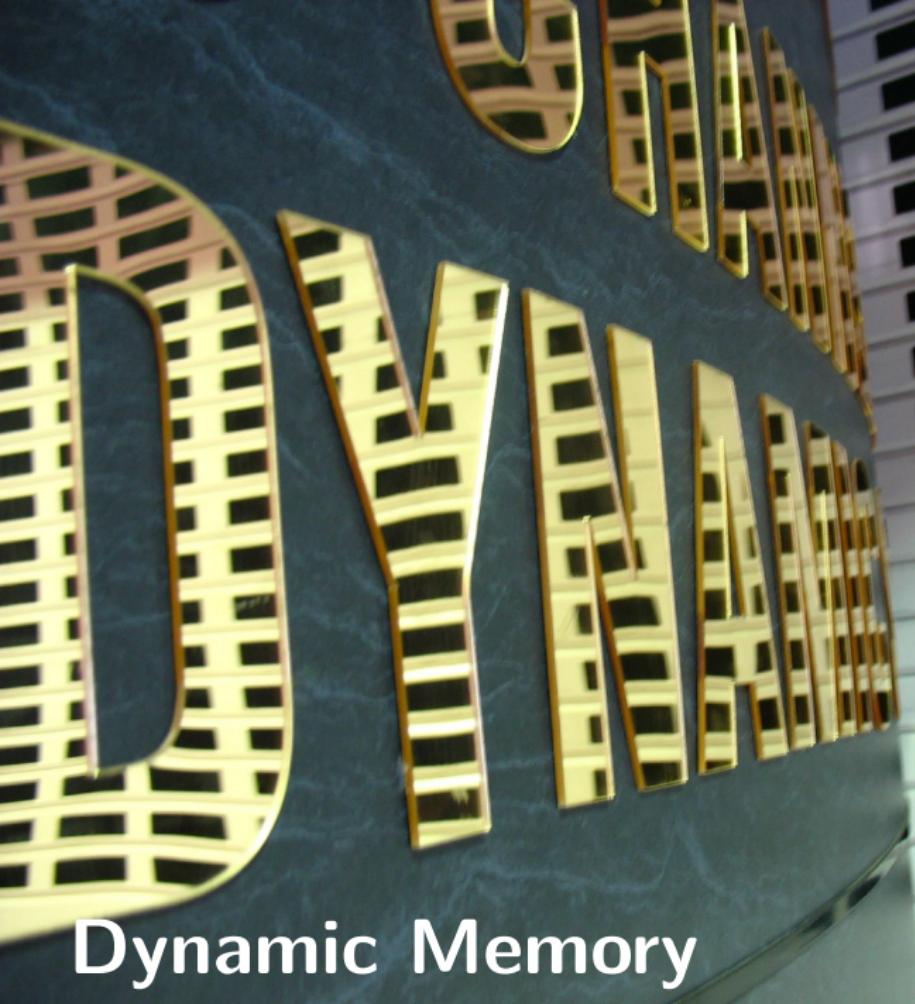
2032	
2024	0
2016	2004
arr:	2008
b:	2004
a:	2000

Dynamic Size Array (2)

- Possible to create manipulation **utility functions**

For example, getting the size of the array by looking for the NULL

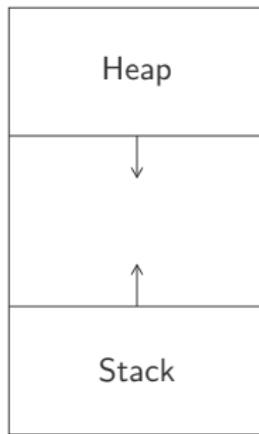
```
1 int length(int *arr[])
2 {
3     int i = 0;
4     while (arr[i++] != NULL);
5     return i;
6 }
7
8 void print_arr(int *arr[])
9 {
10     if (*arr == NULL)
11         printf(" []\n");
12     else {
13         printf(" [%d", **arr);
14
15         while (*(++arr) != NULL)
16             printf(", %d", **arr);
17         printf(" ]\n");
18     }
19 }
```



Dynamic Memory

Dynamic Memory (1)

- There are **two types of memory** in a program
 - The **stack** contains local variables
 - The **heap** contains dynamic variables



Dynamic Memory (2)

- Memory management functions in stdlib.h
 - Creating a dynamic memory area with malloc
 - Freeing the allocated memory area with free

```
1 int *p = malloc(sizeof(int));  
2  
3 printf("%p\n", p);  
4 *p = 12;  
5 printf("%d\n", *p);  
6  
7 free(p);
```

```
0x7fdf58c03970  
12
```

malloc Function

- The `malloc` function can be used to **allocate memory space**
 - Return the start address of the allocated memory block
 - Specifying the desired size in bytes
- The **demand can be refused** if the heap collides with the stack
If unsuccessful, returns a `NULL` pointer

```
1 double *data = malloc(10 * sizeof(double));  
2  
3 if (data != NULL)  
4 {  
5     // ...using the allocated memory area...  
6 }
```

free Function

- The free function can be used to **deallocate memory space**
 - Free space for future memory space allocation demands
 - Does not change the pointer, only invalidate the position
- **Importance of freeing** any memory area not used anymore

Undefined behaviour if freeing a wrong pointer

```
1 double *data = malloc(10 * sizeof(double));
2
3 if (data != NULL) {
4     // ...using the allocated memory area...
5
6     free(data);
7 }
```

Returning Pointer (4)

- The dynamic memory can be used to **share variables**
 - Memory accessible from all the procedures/functions
 - Not deleted after the execution of procedures/functions
- Possibility to write a **function returning a pointer**

Just point to a dynamic memory area

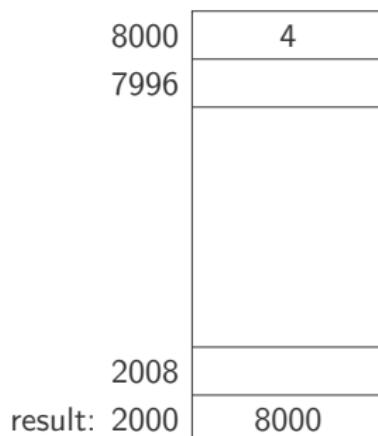
Returning Pointer (5)

```
1 int* new_int(int value)
2 {
3     int *a = malloc(sizeof(int));
4     *a = value;
5     return a;
6 }
```

	8000	4
	7996	
	2028	
a:	2020	8000
value:	2016	4
\$:	2008	2000
result:	2000	

Returning Pointer (6)

```
1 int main()
2 {
3     int *result = new_int(4);
4     printf("%d\n", *result);
5     free(result);
6
7     return 0;
8 }
```



calloc and realloc Functions

- The **calloc function** creates a memory area initialised to 0

```
1 double *data = calloc(5, sizeof(double));
2 if (data != NULL) {
3     // ...using the allocated memory area...
4     free(data);
5 }
```

- The **realloc function** modifies the size of a memory area

```
1 double *data = calloc(5, sizeof(double));
2 if (data != NULL)
3 {
4     data = realloc(data, 10 * sizeof(double));
5     // ...using the allocated memory area...
6     free(data);
7 }
```

Dynamic Array

- Memory allocated by malloc seen as a **dynamic array**

```
1 int N = 3;
2 int *data = malloc(N * sizeof(int));
3
4 int i;
5 for (i = 0; i < N; i++)
6     data[i] = i + 1;
7 print_arr(data, N);
```

8008	3
8004	2
8000	1
7996	
2016	
i:	2012 3
data:	2004 8000
N:	2000 3

Returning Array

- Returning an array must be done with dynamic memory

Do not forget to free the memory after use

```
1 int* new_tab(int size)
2 {
3     return calloc(size, sizeof(int));
4 }
5
6 int main()
7 {
8     int *data = new_tab(3);
9     // ...using the allocated memory area...
10    free(data);
11
12    return 0;
13 }
```

References

- Shohei Yokoyama, *Understanding Memory Layout*, November 10, 2018.
<https://medium.com/@shohei yokoyama/understanding-memory-layout-4ef452c2e709>
- Neotam, *Endianness. Little Endian vs Big Endian*, April 6, 2019.
<https://getkt.com/blog/endianness-little-endian-vs-big-endian>
- Colin Walls, *Endianness*, October 16, 2015.
<https://www.embedded.com/design/mcus-processors-and-socs/4440613/Endianness>
- Puneet Sapra, *No more fear of pointers: Coding Allusion*, February 4, 2018.
<https://medium.com/the-mighty-programmer/variable-and-pointer-fb637566bfd9>

Credits

- Brian Turner, February 21, 2013, <https://www.flickr.com/photos/60588258@N00/8552885919>.
- Colin, September 9, 2018, <https://www.flickr.com/photos/colinsd40/43917455784>.
- Jen, July 27, 2013, <https://www.flickr.com/photos/nstop/9406739132>.
- kishjar?, November 8, 2013, <https://www.flickr.com/photos/kishjar/10747531785>.
- Marc Kjerland, January 12, 2010, <https://www.flickr.com/photos/marckjerland/4285433391>.
- Orin Zebest, May 27, 2006, <https://www.flickr.com/photos/orinrobertjohn/154616556>.