

## Session 2

**Key-Value Model:  
Riak, Memcached, Redis**



This work is licensed under a Creative Commons Attribution – NonCommercial – NoDerivatives 4.0 International License.

# Objectives

- The **key-value** model
  - Principle and characteristics of key-value storage
  - Use case and non-use cases
  - Data repartition models
- **Examples** of key-value databases
  - Riak
  - Memcached
  - Redis

# Key-Value Model



# Key-Value (1)

- Key-value databases similar to hashtables

*Stores key-value pairs, identifiable by their key*

- Similar to a relational table **with two columns**

*Used when searching on primary key*

- Very good performance thanks to **indexing on the key**

<b>Id</b>	<b>Name</b>
16133	Yannis
16067	Théo
16050	Yassine
15089	Maxime

# Key-Value (2)

- The simplest NoSQL storage space

*Regarding the API to use it*

- Mainly three operations on the store

*Retrieve/set a value for a key, delete a key*



LEVELDB

# Data Type

- The stored value is a **blob type** (*Binary Large OBject*)  
*It is up to the application to manage the values and their format*
- Sometimes limits on the **size** of stored values  
*For performance reasons*
- Sometimes **domain constraints** on aggregates  
*Redis supports lists, sets and hashes*

# Basic API

- Three basic operations supported by all engines
  - `get(k)` retrieves the  $v$  value associated to the  $k$  key
  - `put(k, v)` adds the  $(k, v)$  pair in the store
  - `delete(k)` deletes the pair associated to the  $k$  key
- The engine can propose specific operations

*Redis proposes the union of sets, for example*

# Use Case

- Storing **session information** for a website

*Unique identifier convenient for a key-value database*

- **Profiles and preferences** of a given user

*User is characterised by a unique username*

- **Shopping carts** on an e-commerce website

*Storing the current shopping cart of a user*

# Non-Use Case

- Links to establish between data related to different keys

*Following the links between data is not easy*

- Backup of several keys and failure of some backups

*Not possible to restore operations already realised*

- Not possible to make requests on the values

*Except for some specific engines*

# Distribution Model



# Distribution Model

- Several possible models to **operate a cluster**  
*End of scale up (larger server) for scale out (more servers)*
- The **aggregate** information unit can be easily distributed  
*Fine granulometry of information*
- **Several reasons** to use a cluster
  - Ability to manage larger amounts of data
  - Provide a larger read/write traffic
  - Resist to network slowdowns or failures

# Unique Server

- No distribution in the simplest version

*Execution on a single machine that manages reads/writes*

- Solution very simple to implement and operate
  - Easy to manage for operators
  - Easy to reason for application developers
- Suitable for graph-oriented databases

*Where operations to perform are often aggregations*

# Sharding (1)

- Store should be **busy with several users**

*When they are accessing different parts of the data*


- **Sharding** places data on several servers

*Horizontal scalability with deployment of several nodes*

- **Load balancing** between the different servers

*If the users are requesting different data*

# Sharding (2)



# Load Balancing

- Ideally, the **load** is well distributed between clients
  - With 5 nodes, each node manages 20% of the load*
- **Data accessed together** must be placed on the same node
  - Using aggregate as the distribution unit
  - Using the geographical location of data
  - Collecting aggregates by common access probability
- Possibility to have **automatic sharding**
  - The engine manages the sharding and data rebalancing*

# Master-Slave Replication (1)

- Data replicated on several nodes

*Suitable when more reads than writes*


- Two kinds of nodes in the system

- A master node responsible for data and update
- Several slave nodes that are replicates of the master

- Two properties for this kind of replication

- Read resilience allows reads if the master fails
- Values read by users may differ by inconsistency

# Master-Slave Replication (2)



# Data Scattering

- Routing requests based on the type

*Read sent to the slaves and writes to the master*

- Slaves synchronisation by replication process

- Modifications on the master are communicated to the slaves
  - Election of a slave as the master if it fails

- Two modes of choice of the master

- Manual choice by configuration
  - Automatic choice by dynamic election

# Peer-to-Peer Replication (1)

- Data replicated on several nodes that are all equal

*Brings scalability for write operations*


- Synchronising all the nodes at each write

*Concurrent and permanent write conflicts, not like with read*

- Several properties for this kind of replication

- Complete read and write resilience
- Values read by different users different by inconsistency

## Peer-to-Peer Replication (2)



# Sharding vs. Replication

- **Sharding** distributes the load, no resilience

*Different data on different nodes*

- **Replication** offers resilience, heavy synchronisation

*Same data places on different nodes*

Strategy	Scaling	Resilience	Inconsistency
Sharding	Write	–	–
M/S Replication	Read	Read	Yes
P2P Replication	Read/Write	Read/Write	Yes

# Combining Sharding and Replication

- Master-slave replication and sharding
  - Possibility to have several masters, but only one by data
  - Node with a single role or mixed roles
- Peer-to-peer replications and sharding
  - Data sharded on hundreds of nodes
  - Data is replicated on  $N$  nodes (replication factor)



Riak

# Riak

- Created and developed by the **Basho company**

*Company founded in 2008 and develops Riak and other solutions*

- Active company and last version in **may 2019**

*Riak is developed in Erlang and the last version is Riak 2.9.0*

- **Decentralised** NoSQL engine based on Amazon Dynamo

*Scales by adding new machines to the cluster*


# Bucket

- Riak can store keys in **buckets**


*Acts as a namespace for keys*

- **Several possibilities** to operate buckets

*Composed values or separation as “specific objects”*



*versus*



# Domain Bucket

- Domain bucket can store a precise type of data

*Automatic serialisation/deserialisation by the client*

- Separation in buckets to segment data

- Possible to only read objects that you want to read
- Possible to use the same key through different buckets

- Fight against impedance mismatch

*Store directly contains application objects*

# Installing Riak

- Riak is a program written in **Erlang**
- **Several programs** proposed after installation
  - riak to control Riak nodes
  - riak-admin for administration operations

# Starting a Node

- Starting a Riak node with the riak executable

*Starting with the start option and stopping with the stop option*

```
& riak start
```

```
& riak ping  
pong
```

# riak Python Module

- **riak** Python module to query the store

*Opening a connection and then methods to make queries*

```
1 import riak
2
3 client = riak.RiakClient(protocol='http', http_port=8098)
4
5 print(client.ping())
6 print(client.get_buckets())
```

```
True
[]
```

# Creating a Bucket

- Creating a **new bucket** with the `bucket` method

*To be called on the Riak client*

- Return a **RiakBucket** object

*Used to add and read key-value pairs*

```
1 import riak
2
3 client = riak.RiakClient(protocol='http', http_port=8098)
4
5 bucket = client.bucket('students')
6 print(bucket)
```

```
<RiakBucket 'students'>
```

# Data Manipulation

- Creating a **new data** with the **new** method

*Return a RiakObject object that can be stored*

```
1 import riak
2
3 client = riak.RiakClient(protocol='http', http_port=8098)
4 bucket = client.bucket('students')
5
6 print(bucket.get('16050').data)
7
8 yassine = bucket.new('16050', 'Yassine')
9 yassine.store()
10 print(bucket.get('16050').data)
```

```
None
Yassine
```

# Riak Cluster

- Distributing data with a **consistent hash**
  - Minimises keys remapping when the number of nodes changes
  - Distributed the data well and minimises hotspots
- Using **SHA-1** and the 160 bits spaces as ring
  - Cutting the ring in partitions called “virtual nodes”
  - Each physical node hosts several vnodes



KNOXVILLE

Memcached

# Memcached

- General purpose **distributed cache** system

*Speed up a website by caching objects in RAM*

- Used in **combination** with another database

*For example from PHP as a cache to a MySQL database*

- Memcached is a program written in **C**

# Architecture (1)

- Built on a **client/server** architecture

*Server services exposed on the 11211 port by default*

- The client makes **queries by key** on the store

*Keys are at most 250 bytes and values are up to 1 Mio*

- A client knows **all the servers**

- Servers do not communicate between them
- Computation of a hash on the key to chose the server

# Architecture (2)

- Store data are **stored in RAM**
  - Oldest values deleted if not enough RAM
  - Memcached to be used as a transient cache
- Act as a big **hashtable**

*Key-value pairs are stored in this hashtable*

# memcache Python Module

- **memcache Python module** to query the store

*Opening a connection and methods for commands*

```
1 import memcache  
2  
3 mc = memcache.Client(['127.0.0.1:11211'])  
4  
5 print(mc.get('16133'))  
6 print(mc.set('16133', 'Yannis'))  
7 print(mc.get('16133'))  
8 print(mc.delete('16133'))  
9 print(mc.get('16133'))
```

```
None  
True  
Yannis  
1  
None
```

# The Trivago Example

- Trivago uses Memcached for its cache layer

*Avoid a lot of direct requests to the main database*


- Big sudden issue with logs filled with Memcached errors
  - Failures of get and overload of the database
  - Botnet from more than 200 countries with 70K unique IPs...
  - Memcached network interface saturation beyond 1 Gbit/s

# The Facebook Example (1)

- Facebook uses Memcached for a distributed store

*Distributed storage of key-value pairs in memory*


- Two different usages for request or generic
  - Used as a *demand-filled look-aside cache*
  - And also deployment of a generic distributed store



# The Facebook Example (2)

- **No coordination** server-server with Memcached
  - “Only” a local in-memory hashtable of a server*
- **Replication** inside a server cluster

*Data flow from the master to the slaves*





Redis

# Redis

- Database engine **in memory**

*Manipulate data structure as quickly as possible*

- Also plays the role of a **data cache**

*Similar to Memcached with a richer and stronger model*

- Restriction on the **manipulated values**

*Five possible kinds of values stored in the database*

# Value Type

- Possible to manipulate specific **data types** with Redis  
*And do not manipulate documents like other databases*
- **Five** different types of data
  - Strings, and numeric or binary value
  - Lists of strings (insertion order maintained)
  - Set of strings, unsorted and without duplicate
  - Hash (dictionary), not hierarchical
  - Sorted set with association of a note for each element

# Installing Redis

- Redis is a program written in C
- Several programs proposed after installation
  - redis-server to start a Redis server
  - redis-cli is a command-line client
  - redis-benchmark makes a performance test

# Starting the Server

- Starting the server and testing the connection

*Test of a ping to the server from the command line*

```
& redis-server
```

```
& redis-cli  
127.0.0.1:6379> ping  
PONG
```

# Manipulating String

- Several basic commands to **manipulate strings**
  - SET adds a new string in the store
  - GET retrieves the value associated to a key
  - DEL deletes a key from the store

```
& redis-cli
127.0.0.1:6379> GET 15089
(nil)
127.0.0.1:6379> SET 15089 "Maxime"
OK
127.0.0.1:6379> GET 15089
"Maxime"
127.0.0.1:6379> DEL 15089
(integer) 1
127.0.0.1:6379> GET 15089
(nil)
```

# redis Python Module

- **redis Python module** to query the store

*Opening a connection then methods for commands*

```
1 import redis
2
3 r = redis.StrictRedis(host='localhost', port=6379, db=0)
4
5 print(r.get('15089'))
6 print(r.set('15089', 'Maxime'))
7 print(r.get('15089'))
8 print(r.delete('15089'))
9 print(r.get('15089'))
```

```
None
True
b'Maxime'
1
None
```

# Manipulating Hash

- Several basic commands to **manipulate hashes**
  - HSET adds an entry in the hash table of a key
  - HVALS retrieves the complete hash table of a key
  - HGET retrieves the value of an entry of a hash table
  - HDEL deletes an entry of a hash table

```
& redis-cli
127.0.0.1:6379> HSET 16067 firstName Théo
(integer) 1
127.0.0.1:6379> HSET 16067 favColour green
(integer) 1
127.0.0.1:6379> HVALS 16067
1) "Théo"
2) "green"
127.0.0.1:6379> HGET 16067 favColour
"green"
```

# Hash/Python Dictionary Equivalence

- Direct mapping between hashes and **Python dictionaries**

*Initialisation of a hash with hmset*

```
1 import redis
2
3 r = redis.StrictRedis(host='localhost', port=6379, db=0)
4 r.hmset('10003', {
5     'firstName': 'Théo',
6     'favColour': 'green'
7 })
8 print(r.dbsize())
9 print(r.hgetall('10003'))
```

```
1
{b'firstName': b'Théo', b'favColour': b'green'}
```

# Manipulating List

- Several basic commands to **manipulate lists**
  - LPUSH adds an entry to the left of a list
  - LPOP removes the entry to the left of a list
  - RPUSH adds an entry to the right of a list
  - RPOP removes the entry to the right of a list
  - LRANGE extract a sublist from a list

```
& redis-cli
127.0.0.1:6379> RPUSH students 16133
(integer) 1
127.0.0.1:6379> RPUSH students 15089
(integer) 2
127.0.0.1:6379> LRANGE students 0 -1
1) "16133"
2) "15089"
```

# List/Python List Equivalence

- Direct mapping between lists and **Python lists**

*Initialisation of a list with rpush*

```
1 import redis
2
3 data = [ '16133', '15089' ]
4
5 r = redis.StrictRedis(host='localhost', port=6379, db=0)
6 r.delete('students')
7 r.rpush('students', *data)
8
9 data = r.lrange('students', 0, -1)
10 for elem in data:
11     print(elem)
```

```
b'16133'
b'15089'
```

# Data Persistence

- Redis is a **in-memory only** database  
*Once the server exits, all data is lost*
- Possibility to **regularly save** data on disk  
*Using the RDB system by default, for regular snapshots*
- **Automatic reloading** of the database  
*If a .rdb file is in the right folder*

# Expiration

- Possible to choose the **lifetime** of elements

*Using the EXPIRE command*

- An element in a cache should **not live forever**

# Redis Social Network Example

- Storing a simple **social network** with Redis

*Defining the format of key-value pairs to use*

- **Two kinds of objects** in the store

- **User** has a name and can be followed by others
  - **Post** is a message, a picture...

- A user can have **several posts**

*Storing the list of posts of a user*

# Key Format (1)

- Defining the **format of the keys** to use

*Must be a simple string*

- **Convention** to have unique keys

- **User**

```
user:1:name → Mathias  
username:Mathias → 1
```

- **Post**

```
post:1:content → Hi Théo, you rock!  
post:1:user → 1
```

# Key Format (2)

- Posts and follow relations with **lists/sets**

*Integer numbers lists referring users and posts*

- Using “**sub-keys**” from user

- **Posts list**

`user:1:posts → [3, 2, 1]`

- **Follow relation**

`user:1:follows → {2, 3, 4}`

`user:1:followed_by → {3}`

# Automatic Identifier

- Possibility to **increment a value** with the INCR command

*The value must represent an integer number*

- Adding two pairs to represent the **next IDs**

*Keys next\_user\_id and next\_post\_id*

```
1 import redis
2
3 r = redis.StrictRedis(host='localhost', port=6379, db=0)
4 r.set('next_user_id', 0)
5 print(r.get('next_user_id'))
6
7 r.incr('next_user_id')
8 print(r.get('next_user_id'))
```

```
b'0'
b'1'
```

# Creating a New User

- Definition of a method to **create a new user**

```
1 import redis
2
3 r = redis.StrictRedis(host='localhost', port=6379, db=0)
4 r.set('next_user_id', 0)
5
6 def create_user(username):
7     uid = int(r.get('next_user_id'))
8     r.set('user:{}:name'.format(uid), username)
9     r.set('username:{}'.format(username), uid)
10    r.incr('next_user_id')
11
12 create_user('Mathias')
13 create_user('Théo')
14
15 print(r.get('user:0:name'))
16 print(r.get('user:1:name'))
```

```
b'Mathias'
b'Théo'
```

# Top 5 Redis Use Cases

- Session **cache** and Full Page Cache (FPC)

*The advantage of Redis is persistance*

- Implementation of an efficient **message queue**

*For example with the Celery tool for Distributed Task Queue*

- Developing a **leaderboard** with counting

- Execution of scripts with **Pub/Sub events**

# References

- Wishmitha S. Mendis, *From RDBMS to Key-Value Store: Data Modeling Techniques*, October 29, 2017. <https://medium.com/@wishmithasmendis/from-rdbms-to-key-value-store-data-modeling-techniques-a2874906bc46>
- Darren Perucci, DZone, *Redis Replication vs Sharding*, June 15, 2016. <https://dzone.com/articles/redis-replication-vs-sharding>
- Ivana Petrovic and Polina Pokalyukhina, *How trivago Reduced Memcached Memory Usage by 50%*, December 19, 2017. <https://tech.trivago.com/2017/12/19/how-trivago-reduced-memcached-memory-usage-by-50>
- Rajesh Nishtala, Hans Fugal, Steven Grimm, Marc Kwiatkowski, Herman Lee, Harry C. Li, Ryan McElroy, Mike Paleczny, Daniel Peek, Paul Saab, David Stafford, Tony Tung and Venkateshwaran Venkataramani (2013). Scaling Memcache at Facebook. In *Proceedings of the 10th USENIX Symposium on Networked Systems Design and Implementation (NSDI 2013)*. [https://www.usenix.org/system/files/conference/nsdi13/nsdi13-final170\\_update.pdf](https://www.usenix.org/system/files/conference/nsdi13/nsdi13-final170_update.pdf)
- Joe Engel, *Top 5 Redis Use Cases*, November 7, 2017. <https://www.objectrocket.com/blog/how-to/top-5-redis-use-cases>

# Credits

- Logo pictures from Wikipedia.
- SioW, July 3, 2006, <https://www.flickr.com/photos/curioussiow/182224885>.
- Shepherd Distribution Services, October 15, 2010,  
<https://www.flickr.com/photos/shepherd-distribution-services/5395849861>.
- <https://openclipart.org/detail/94723/database-symbol>.
- heschong, May 14, 2007, <https://www.flickr.com/photos/heschong/510216272>.
- DM, April 27, 2011, <https://www.flickr.com/photos/dmott9/5662744650>.
- othreee, November 19, 2013, <https://www.flickr.com/photos/othreee/10945272436>.