

## Session 3

# Column-Oriented Model: Cassandra, HBase



This work is licensed under a Creative Commons Attribution – NonCommercial – NoDerivatives 4.0 International License.

# Objectives

- **Column-Oriented** model
  - Storing rows or columns on disk
  - The data model
  - Main types of queries
- **Examples** of column-oriented databases
  - HBase
  - Cassandra

# Column-Oriented Model



# Column Family (1)

- Column-oriented databases close to relational ones

*Include columns with a given data type*

- Follow the BigTable approach brought by Google

*Whose HBase is an open source implementation*

- Quick access to data and very good scalability

*In particular with Cassandra and a peer-to-peer distribution*

# Column Family (2)

- Set of **row keys** and column families

*Organisation of a database with several tables*

- Grouping together data often **accessed together**

*Each column family is a data map*



cassandra



HYPERTABLE INC



druid



# Row vs. Column (1)

- Disk storage **by tuples or by rows**

*Initially only a storage issue*

- Queries do not often include **all columns**

*Direct column retrieval from the disk more efficient*

ID	Firstname	Class
16067	Théo	4MIN
15056	Houda	5MIN

*Stockage de lignes*

ID	Firstname	Class
16067	Théo	4MIN
15056	Houda	5MIN

*Stockage de colonnes*

# Row vs. Column (2)

- Choosing the disk storage to have **efficient operations**
  - Row storage efficient for writes
  - Row storage efficient for reads
- Reading of a **few columns** with many rows

*Improve the performances of select queries*

Row storage	Column storage
+ Easy to add a record	Only the desired data is read
- Reading unnecessary data	Writing a tuple requires multiple access

# C-Store (1)

- Storing the data **in columns** in the database

*Created by Brown, Brandeis, MIT and UMass Boston universities*

- Based on the **relational model** and uses SQL

*Does not belong to the NoSQL world, but will inspire it*

- **Two different storage spaces** on the disk

*To better optimise the read and write operations*

# C-Store (2)

- **ROS** (*Read Optimized Store*)
  - Storing files containing columns
  - Compressing files depending on the included data types
  - Data sorted by an attribute of the table or the column
- **WOS** (*Write Optimized Store*)
  - Temporary buffer used for write (INSERT, UPDATE)
  - No compression and vertical partitioning

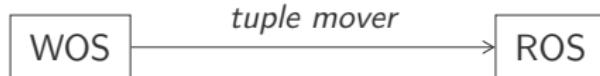
# C-Store (3)

- Regular **migration** of data from the WOS to the ROS

*Realised by a tuple mover authorised to write in the ROS*

- **Queries** must be able to operate on both stores

- Insertions directly sent to the WOS
- Deletions marked in the ROS, then managed by tuple mover
- Update is a combination of insertions and deletions



# Row vs. Column (3)

- No absolute best choice between rows and columns

*It depends on the kind of performed operations*

	Rows	Columns
Aggregating elements from a column	Slow	Fast
Compression	–	High
Selecting a few columns	Slow (skipping data)	Fast
Insertion/Update	Fast	Slow
Selecting a record	Fast	Slow

# Data Model (1)

- A column-oriented base is a **two-level map**

*Rather than a table structure organised by columns*

- A **key-value pair** identifies a row at the first level

*The key is a row identifier*

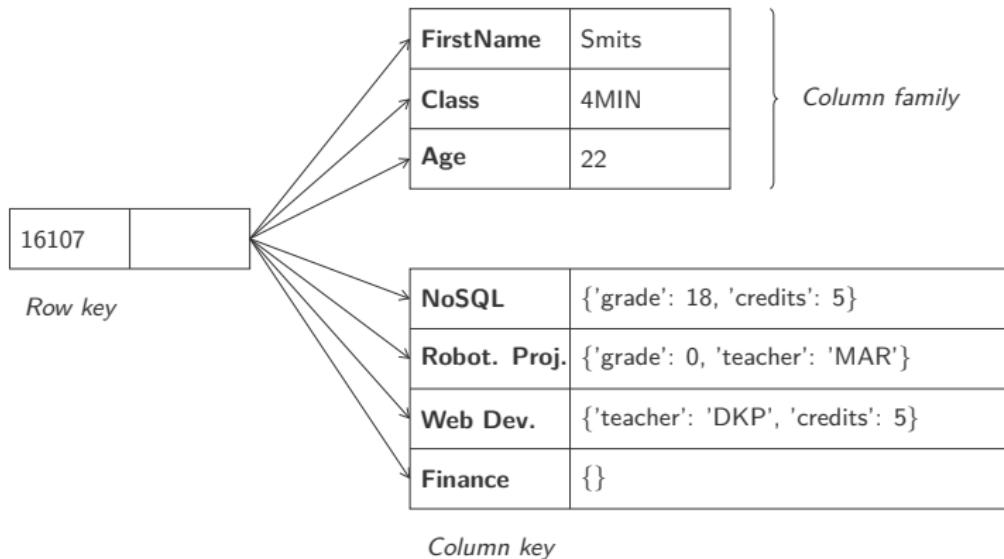
- A **map of columns** forming families at the second level

- Arbitrary number of key-value pairs by row
- Families for common accesses to columns

# Data Model (2)

- Two-level structure combining **rows** and **columns**

*Row is the join of records from column families*



# Data Model (3)

- Column-oriented databases are **not really tables**
  - Columns can be added to any row
  - Rows can have different column keys
- Defining **new column families** is rare

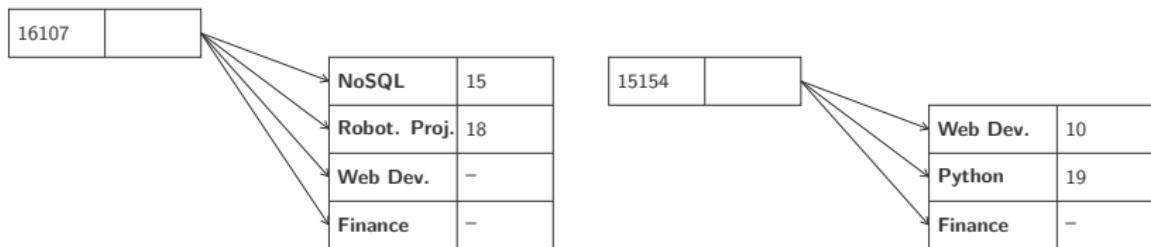
*But adding new column can be done on the fly*
- **Two kinds** of rows depending on the number of columns
  - **Skinny row** few columns and same everywhere (*field-like*)
  - **Wide row** thousands of columns (*list-like*)

# Table vs. Column

- Column-oriented databases avoid **presence of NULL**

*Each row only has the columns it should have*

Matricule	NoSQL	Robot. Proj.	Web Dev.	Python	Finance
16107	18	0	–	NULL	–
15154	NULL	NULL	10	19	–



# Column Advantage

- Efficient read of data only from the necessary columns

*Watch out for tuple reconstruction when reading all*

- Better compression rate, but higher CPU usage

*Less entropy since all data from the same domain*

- Efficiency of data sorting and indexing

*With redundant storage thanks to space gained by compression*

# Projection (1)

- Possibility to have physically stored **projections**

*To improve performances for some query types*

**Logical table**

Region	Customer	Product	Sale
A	G	C	789
B	C	C	743
D	F	D	675
C	C	A	23
A	R	B	654

**Super-projection**

Region	A	B	D	C	A
Customer	G	C	F	C	R
Product	C	C	D	A	B
Sale	789	743	675	23	654

# Projection (2)

- Projections can be **sorted** on one or several columns

*Improve performance for *SORT* and *GROUP BY* requests*

**Logical table**

Region	Customer	Product	Sale
A	G	C	789
B	C	C	743
D	F	D	675
C	C	A	23
A	R	B	654

**Projection 1**

Region	A	A	B	C	D
Product	B	C	C	A	D
Sale	654	789	743	23	675

Ease query such as:

```
SELECT Region, Product, SUM(Sale)  
GROUP BY Region, Product
```

# Projection (3)

- Can be created manually or **on the fly**

*A bit the same logic than having materialised views*

**Logical table**

Region	Customer	Product	Sale
A	G	C	789
B	C	C	743
D	F	D	675
C	C	A	23
A	R	B	654

**Projection 2**

Customer	C	C	F	G	R
Sale	743	23	675	789	654

Ease query such as:

```
SELECT Customer, SUM(Sale)  
GROUP BY Customer
```

# Compression (1)

- Run-Length Encoding on values in the columns

*Convenient when a lot of similar data*

Semester	Product	Price	Semester	Product	Price
Q1	1	5	(Q1, 1, 300)	(1, 1, 4)	5
Q1	1	7	(Q2, 301, 350)	(2, 5, 2)	7
Q1	1	2	...	...	2
Q1	1	9			9
Q1	2	6			6
Q1	2	8			8
...	...	...			...
Q2	1	3			3
Q2	1	8			8
Q2	2	1			1
...	...	...			...

# Compression (2)

- Bit-Vector Encoding for each unique value of columns

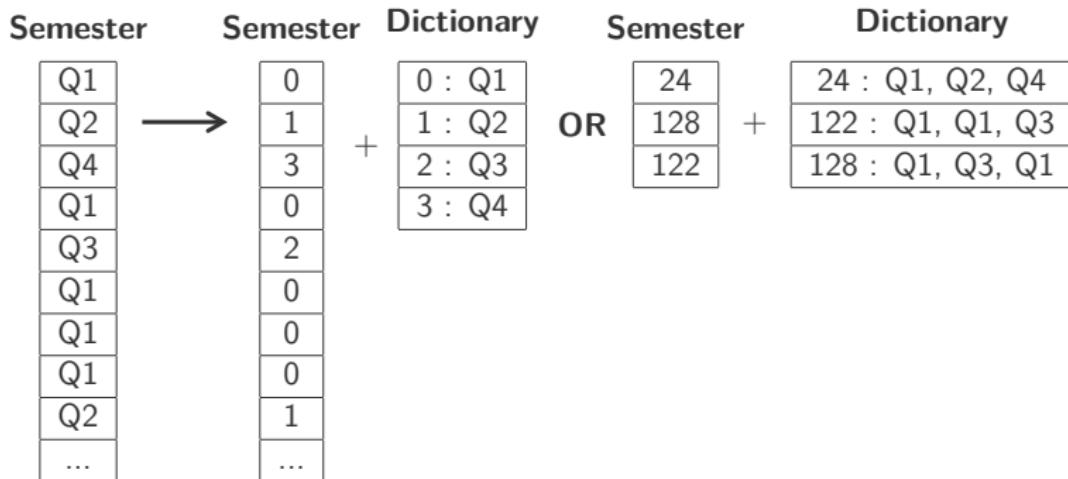
*Convenient when only few unique values, combined with RLE*

Product	Product : 1	Product : 2
1	1	0
1	1	0
1	1	0
1	1	0
2	0	1
2	0	1
...	...	...
1	1	0
1	1	0
2	0	1
...	...	...

# Compression (3)

- Dictionary for each value or block of values

*Convenient when pattern repetitions*



# Use Case

- Storing **events logs**

*State changes or errors found in an application*

- **Blog posts** as part of a CMS

*Tags, categories, links, etc. in different columns of a family*

- **Count and categorise** visitors of a webpage

*Using a particular counter type column*

# Non-Use Case

- Problems for which **ACID must be satisfied** for read/write

*No ACID transactions with column-oriented databases*

- **Data aggregation** requests (SUM, AVG, etc.)

*First requires to get all the rows on the client side*

- Do not use when in a **prototyping phase**

*The design of column families change with requests to perform*

'One of the greatest thinkers of the age' the Dalai Lama  
J. KRI SHNAMURTI



HBase

FREEDOM *from the* KNOWN

THINK ON THESE THINGS

J KRIS

Programming Pig

HBase *The Definitive Guide*

# HBase

- Open source implementation of the **BigTable** engine by Google

*Is part of the Hadoop project by Apache*

- Executed on top of the **HDFS** file system

*Storage of sparse data while being fault-tolerant*

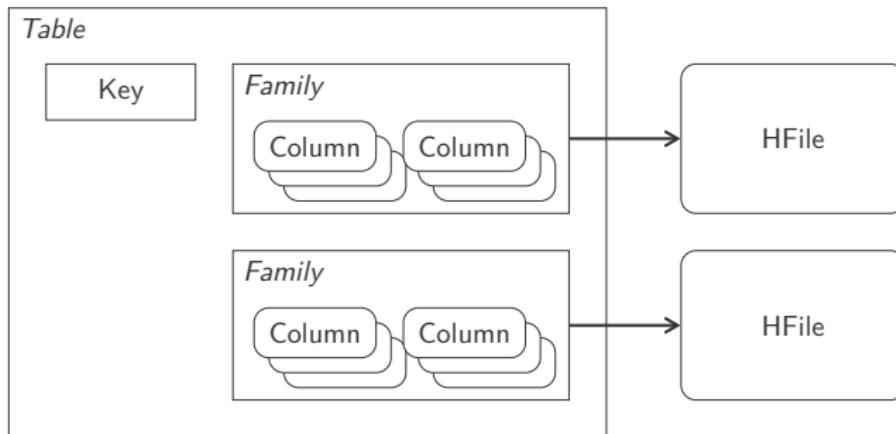
- A DB can serve as **input/output of MapReduce** (Hadoop)

*Possible to have a SQL layer thanks to Apache Phoenix*

# Data Model

- Set of versioned **column families**

*Columns of a given family stored together in a HFile*



Path to find a value: Table → Key → Family → Column → Timestamp

# Architecture (1)

- Based on Hadoop and HDFS to **distribute the storage**

*Combination of sharding and replication*

- Sharding realised by **region servers**

*Split in several regions when a table becomes too big*

- Replication ensured **automatically** by HDFS

*File split in blocks replicated with a given factor*

# Architecture (2)

- Written data are going through several steps
  - First handled in a WAL (*Write-Ahead Log*)
  - Data places in a buffer named *memstore*
- Memstore writes in a **HFile on the HDFS** when too big

*Sorted set of key-values serialised on disk and immutable*
- Deletion managed thank to a *tombstone* marker

*Effective deletion at the same time than compaction*

# Installing HBase

- HBase is a program written in **Java**
- **Several programs** proposed after installation
  - `start-hbase` is a script that starts an HBase server
  - `stop-hbase` is a script that stops an HBase server
  - `hbase` is used to launch several management commands
    - `hbase shell` proposes a command line interface client
    - `hbase thrift` starts the Thrift gateway

# Starting the Server

- Starting the server and verifying the connection

*Using status to check that everything is good*

```
& start-hbase.sh
```

```
& hbase shell
HBase Shell; enter 'help<RETURN>' for list of supported commands.
Type "exit<RETURN>" to leave the HBase Shell
Version 1.2.2, r3f671c1ead70d249ea4598f1bbcc5151322b3a13, Fri Jul
1 08:28:55 CDT 2016

hbase(main):001:0> status
1 active master, 0 backup masters, 1 servers, 0 dead, 2.0000
average load
```

# Creating a Table

- Creating a **new table** with the `create` command

*Specifying column families with the number of versions*

```
hbase(main):002:0> create 'students', {NAME => 'infos', VERSIONS => 1}, {NAME => 'registrations', VERSIONS => 2}
0 row(s) in 1.2230 seconds

=> Hbase::Table - students

hbase(main):003:0> list
TABLE
students
1 row(s) in 0.0630 seconds

=> ["students"]
```

# Adding a Row

- Adding **values to different columns** with `put`

*Specifying each time the column family*

```
hbase(main):004:0> put 'students', '16107', 'infos:firstname', 'Smits'
0 row(s) in 0.1350 seconds

hbase(main):005:0> put 'students', '16107', 'infos:age', '22'
0 row(s) in 0.0120 seconds

hbase(main):006:0> put 'students', '16107', 'registrations:class',
  ', '4MIN'
0 row(s) in 0.0110 seconds

hbase(main):007:0> get 'students', '16107'
COLUMN                           CELL
infos:age                         timestamp=1477172359150, value=22
infos:firstname                     timestamp=1477172339414, value=Smits
registrations:class               timestamp=1477172463762, value=4MIN
3 row(s) in 0.0750 seconds
```

# New Version of a Column

- Possible to retrieve the **different versions of a column**

*Using parameters of the get command*

```
hbase(main):008:0> put 'students', '16107', 'registrations:note',
  'Loves electronics'
0 row(s) in 0.0030 seconds

hbase(main):009:0> put 'students', '16107', 'registrations:note',
  'Loves informatics'
0 row(s) in 0.0030 seconds

hbase(main):010:0> get 'students', '16107', {COLUMN => 'registrations:note', VERSIONS => 2}
COLUMN          CELL
  registrations:note      timestamp=1477173105470, value=Loves
  informatics
  registrations:note      timestamp=1477173102196, value=Loves
  electronics
2 row(s) in 0.0110 seconds
```

# happybase Python module

- happybase Python module to query the database

*Thrift gateway to start with hbase thrift start*

```
1 import happybase
2
3 connection = happybase.Connection('localhost')
4 print(connection.tables())
5
6 table = connection.table('students')
7 print(table)
```

```
[b'students']
<happybase.table.Table name=b'students'>
```

# Inserting a Column

- Columns insertion with the `put` method of the table

*The different columns are provided by a dictionary*

- Row columns retrieval with the `row` method

```
1 table.put('15154', {  
2     'infos:firstname': 'Mathias',  
3     'infos:sex': 'M',  
4     'registrations:class': '4MIN'  
5 })  
6 print(table.row('15154'))
```

```
{b'infos:sex': b'M', b'infos:firstname': b'Mathias', b'  
registrations:class': b'4MIN'}
```

# Retrieving Columns

## ■ Retrieving a row with `row` and several with `rows`

*Possible to filter the columns to only keep the desired ones*

```
1 users = [b'16107', b'15154']
2 classes = {}
3 rows = table.rows(users, columns=[b'infos:firstname', b'
4 registrations:class'])
5 for key, value in rows:
6     students = classes.setdefault(value[b'registrations:class'],
7         set())
7     students.add(value[b'infos:firstname'])
```

```
{b'4MIN': {b'Mathias'}, b'4MIN': {b'Smits'}}
```



Cassandra

# Cassandra

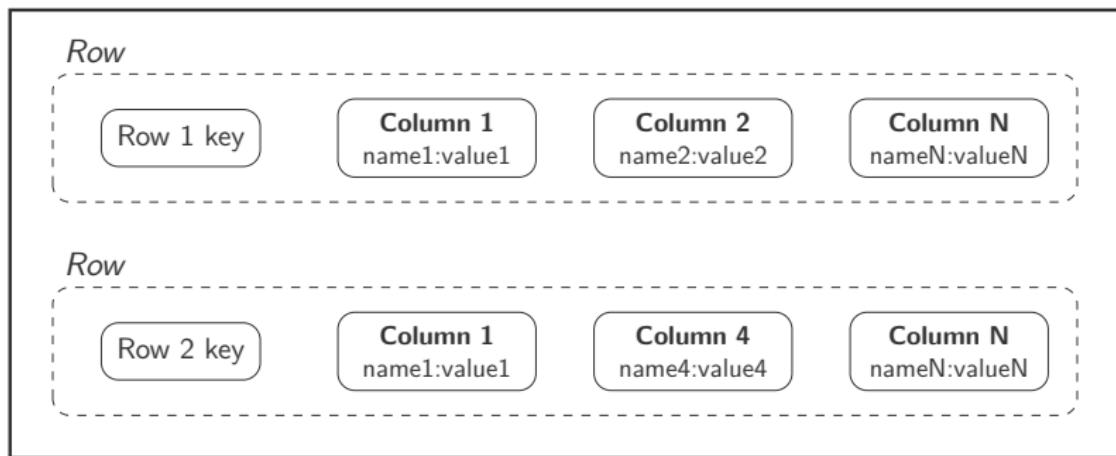
- Originally developed by Facebook and **open sourced in 2008**  
*Is not part of Apache's lap*
- **Fast and scalable** database, peer-to-peer replication on cluster  
*Commodity servers, no single point of failure*
- Query language **Cassandra Query Language** (CQL)  
*Variant of SQL to query Cassandra keyspaces*

# Data Model

- **Column families** set with rows

*Rows can contain different columns of the family*

*Column family*



# Column

- A column is a **key-value pair** with a timestamp

*The name of the column also plays the role of a key*

- The **timestamp** defines the lifetime of the column

*And write conflict resolution, stale data, etc.*

```
1 {  
2   name: "FirstName",  
3   value: "Smits",  
4   timestamp: 1234567890  
5 }
```

```
1 {  
2   name: "Class",  
3   value: "4MIN",  
4   timestamp: 1234567890  
5 }
```

# Standard Column Family

- A **row** is a collection of columns

*A key is attached to this collection of columns*

- A **column family** is a collection of similar rows

*Columns are simple, just a name and a value*

```
1 {  
2     smits: {                      # row with 3 columns, key "smits"  
3         FirstName: "Smits",  
4         Class: "4MIN",  
5         Age: 22  
6     },  
7     mathias: {                      # row with 3 columns, key "mathias"  
8         FirstName: "Mathias",  
9         Class: "4MIN",  
10        Sex: "M"  
11    }  
12 }
```

# Supercolumn

- The value of a **supercolumn** is a map

*“Several columns” as the value of a column*

- A supercolumn is a **container of columns**

*Each contained column has a timestamp*

```
1  {
2      name: "04020",
3      value: {
4          name: "Data acquisition and treatment",
5          coordinator: "MCH",
6          credits: 4
7      },
8      timestamp: 1234567890
9  }
```

# Supercolumn Family

- A **supercolumn family** gathers supercolumns

*Watch out that Cassandra retrieves all, not always optimal*

```
1  {
2      3BE: {
3          E3050: {
4              name: "Signals, systems and telecommunications",
5              coordinator: "DBR",
6              credits: 6
7          },
8          E3010: {
9              name: "Microcontroller and Logic Design",
10             coordinator: "FLE",
11             credits: 6
12         }
13     },
14     4MIN: {
15         04020: {
16             name: "Data acquisition and treatment",
17             credits: 4
18         }
19     }
20 }
```

# Keyspace

- Cassandra organises the column families into **keyspaces**

*Acts like a namespace for column families*

- Similar to the notion of **base** of relational engines

*Gathering families linked to a same application*

# Installing Cassandre

- Cassandra is a program written **in Java**
- **Several programs** proposed after installation
  - `cassandra` starts a Cassandra server
  - `cqlsh` is a client command line interface
  - `nodetool` gives information about Cassandra server

# Starting the Server

- Starting the server and checking the connection

*Immediate indication of whether a server has been found*

```
& cassandra
```

```
& cqlsh
Connected to Test Cluster at localhost:9042.
[cqlsh 5.0.1 | Cassandra 3.7 | CQL spec 3.4.2 | Native protocol
v4]
Use HELP for help.
cqlsh>
```

# Executing a Query

- Obtaining **information on the cluster** with a CQL query

*Information retrieved from the system.local table*

- Great **similarity** with SQL queries

```
cqlsh> SELECT cluster_name, listen_address FROM system.local;  
  
cluster_name | listen_address  
-----+-----  
Test Cluster | 127.0.0.1  
  
(1 rows)
```

# Information on the Base

- Obtaining **information** with the DESCRIBE command

*Description of cluster, keyspaces, tables, etc.*

```
cqlsh> DESCRIBE CLUSTER;

Cluster: Test Cluster
Partitioner: Murmur3Partitioner

cqlsh> DESCRIBE KEYSPACES;

system_traces    system_schema    system_auth    system
system_distributed

cqlsh> DESCRIBE TABLES;

Keyspace system_traces
-----
events    sessions
[...]
```

# Creating a Keyspace

- Creating a **new keyspace** with CREATE KEYSPACE

*Configuring the keyspace properties, for example replication*

- Example with **simple replication** with a given factor

```
cqlsh> CREATE KEYSPACE myschool
... WITH replication={'class': 'SimpleStrategy', '
replication_factor': 3};

cqlsh> DESCRIBE keyspaces;

myschool    system_schema    system_auth    system    system_distributed
system_traces

cqlsh> USE myschool;
cqlsh:myschool>
```

# Creating a Table

- Creating a **new table** with CREATE TABLE

*Definition of the different columns of the table*

- **Primary key** to uniquely identify rows

```
cqlsh:myschool> CREATE TABLE students (
...    serial int PRIMARY KEY,
...    firstName text,
...    class text,
...    age int,
...    sesque text
... );

cqlsh:myschool> SELECT * FROM students;

serial | age | class | firstName | sesque
-----+----+-----+-----+-----
(0 rows)
```

# Adding and Removing Column

- The **table structure** can be changed with ALTER TABLE  
*Possibility to add and remove columns*
- Example of a **correction of the column** sesque in sex

```
cqlsh:myschool> ALTER TABLE students DROP sesque;
cqlsh:myschool> ALTER TABLE students ADD sex text;
cqlsh:myschool> SELECT * FROM students;
    serial | age | class | firstName | sex
-----+-----+-----+-----+
(0 rows)
```

# Adding Row

- Adding a row in the table with `INSERT INTO`  
*Specifying the columns for which there is a value to set*
- Example of adding Smits in the students table

```
cqlsh:myschool> INSERT INTO students (serial, firstName, class,
age)
... VALUES (16107, 'Smits', '4MIN', 22);
cqlsh:myschool> SELECT * FROM students;

  serial | age | class | firstName | sex
-----+-----+-----+-----+-----
  16107 | 22 | 4MIN |      Smits | null
(1 rows)
```

# Other CRUD Operations

- Three other **CRUD operations** as with SQL
  - **Update** rows

```
UPDATE table SET n1=v1, n2=v2... WHERE cond
```
  - **Read** rows

```
SELECT c1, c2... FROM table WHERE cond
```
  - **Delete** rows

```
DELETE c1, c2... FROM table WHERE cond
```
- Operation **on a single row** with a condition on its key

*Not specifying c1, c2... acts on a whole column*

# cassandra Python Module

- **cassandra Python Module** to query the database

*Creation of a cluster and connection on a keyspace*

```
1 from cassandra.cluster import Cluster  
2  
3 cluster = Cluster(['127.0.0.1'])  
4 session = cluster.connect('myschool')  
5  
6 print(cluster)  
7 print(session)
```

```
<cassandra.cluster.Cluster object at 0x1096af240>  
<cassandra.cluster.Session object at 0x10a6bed30>
```

# Executing a Query

- Using the `execute` method on the session

*Executing a CQL query, retrieving a named tupled*

- The class column will **not be accessible** as a field

*Because of a conflict with the `class` property of Python*

```
1 rows = session.execute('SELECT * FROM students')
2 for row in rows:
3     print(row)
4     print('=> {} ({} y.o.)'.format(row.firstName, row.age))
```

```
Row(serial=16107, age=22, field_2_='4MIN', firstName='Smits', sex=None)
=> Smits (22 y.o.)
```

# Building a Query

- Query by **inserting values** in a string

*Similar to formatted outputs*

```
1 session.execute(  
2     '''  
3         INSERT INTO students (serial, firstName, class, sex)  
4             VALUES (%s, %s, %s, %s)  
5         ''',  
6         (15154, 'Mathias', '4MIN', 'M')  
7     )  
8  
9 rows = session.execute('SELECT * FROM students')  
10 for row in rows:  
11     print(row)
```

```
Row(serial=15154, age=None, field_2_='4MIN', firstName='Mathias',  
sex='M')  
Row(serial=16107, age=22, field_2_='4MIN', firstName='Smits', sex=  
None)
```

# Prepared Query

- Building a **prepared query** with the `prepare` method

*Then execution with the `execute` method*

- Authorise search on a column with `ALLOW FILTERING`

```
1  search_class = session.prepare('SELECT class FROM students WHERE
2    firstName=? ALLOW FILTERING')
3
4  users = ['Harold', 'Smits', 'Théo', 'Mathias']
5  classes = []
6  for user in users:
7      rows = session.execute(search_class, [user])
8      for row in rows:
9          students = classes.setdefault(row[0], set())
10         students.add(user)
11
12 print(classes)
```

```
{'4MIN': {'Smits', 'Mathias'}}
```

# References

- Mangat Rai Modi, *Rowise vs Columnar Database? Theory and in Practice*, January 26, 2018.  
<https://medium.com/@mangatmodi/rowise-vs-columnar-database-theory-and-in-practice-53f54c8f6505>
- Ameya, *C-Store: A Columnar Database: Introduction*, April 5, 2019.  
[https://medium.com/@ameya\\_s/c-store-a-columnar-database-1fe7e84d7247](https://medium.com/@ameya_s/c-store-a-columnar-database-1fe7e84d7247)

# Credits

- Logo pictures from Wikipedia.
- zolakoma, August 8, 2008, <https://www.flickr.com/photos/zolakoma/2847597889>.
- balu, May 15, 2014, <https://www.flickr.com/photos/balusss/14004726607>.
- Simon Winch, August 23, 2010, <https://www.flickr.com/photos/110777427@N06/14184365994>.