

Session 1

Definition and Structure of Operating Systems



This work is licensed under a Creative Commons Attribution – NonCommercial – NoDerivatives 4.0 International License.

Objectives

- Define computing system and **operating system** (OS)

What it is, what it does and the relation between both

- Understand the **services provided** by an OS

To the user and to the system through system calls

- Discover how the OS handles **CPU, memory and input/output**

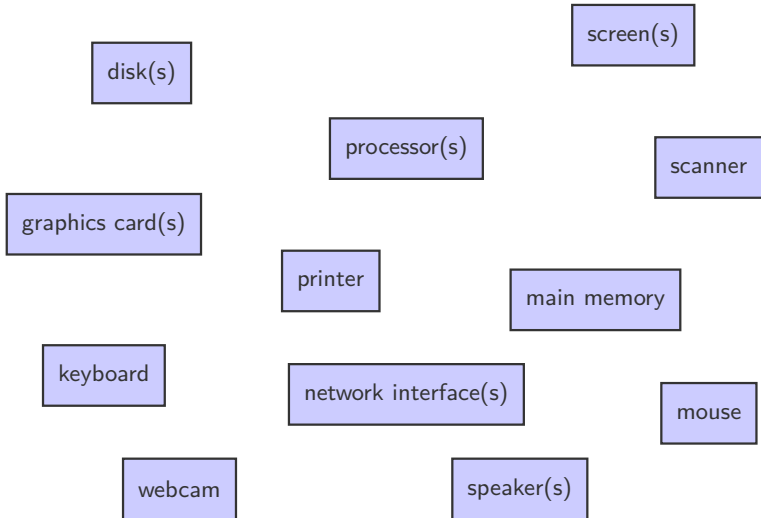
With software abstractions stored and managed by the OS

Operating System


FLATRON 2240

Missing operating system_

Modern Computing System



Modern Computing System



work interface(s)



Computing vs Operating System

- **Computing system** composed of three elements
 - Hardware
 - Software
 - Data
- The OS provides an **environment** to use these resources

No useful functions in itself, used by other programs
- Modern OS **controlled by interruptions**

OS code only executed following events reported by interruptions

Operating System Example



ANDROID



ubuntu



Google Chrome OS

macOS



iOS

UNIX[®]

Celebrating 40 years uptime



Operating System

- Intermediate software layer **between user and hardware**

Provides user programs an interface to the hardware

- **Three main objectives** for an OS

- **Practical**: easy to use by users (PC, mobile)
- **Efficient**: optimise the use of hardware (mainframes)
- **Scalable** : addition of new system functions


- Very large and **very complex** software system

Created piece by piece with a clear delimitation

Computing System Structure (1)

- A **computing system** can be seen with four layers

Each layer uses services from below to provide services to above



Computing System Structure (2)

- The **hardware** provides computing resources

Processor (CPU), memory, input/output devices (I/O)

- The **operating system** coordinates the use of hardware

Coordination between applications and users

- The **applications** solve problems from the users

Word processing, compiler, video game, web browser, etc.

- The **users** uses the computing system

People, machine, other computers, etc.

Definition

- There is no universally adopted **definition**

Great diversity: toaster, car, space shuttle, house, game machine, industrial control system, etc.

- “Everything in the box when you buy an OS”

Microsoft 1998 lawsuit with Internet Explorer

- “Only program running on the computer all the time”

Kernel, system programs and application programs

Operating System Services

S E R V I C E

OS Abstraction

- **Environment** for applications to work

Providing an access to the hardware for applications

- Several **abstractions** are proposed by the OS

Transparent handling of the hardware

Hardware	Abstraction
CPU and memory	Processes and threads
Storage on disk	Files
Network interface	Sockets, RPC, RMI
Display	Drawing libraries and windows

OS Function

- The OS acts as a **service provider**

File system, standard libraries, window system, etc.

- The OS works as a **coordinator** on three aspects
 - **Protection**: jobs must not interfere with each other
 - **Communication**: jobs must be able to interact together
 - **Resource management**: resource sharing must be facilitated


OS Operation

- The **OS does nothing** as long as there is nothing to do

Processes to execute, I/O to handle, users to respond to

- Alternation between two **modes of execution**


Distinguishing the execution of the OS code from the user code



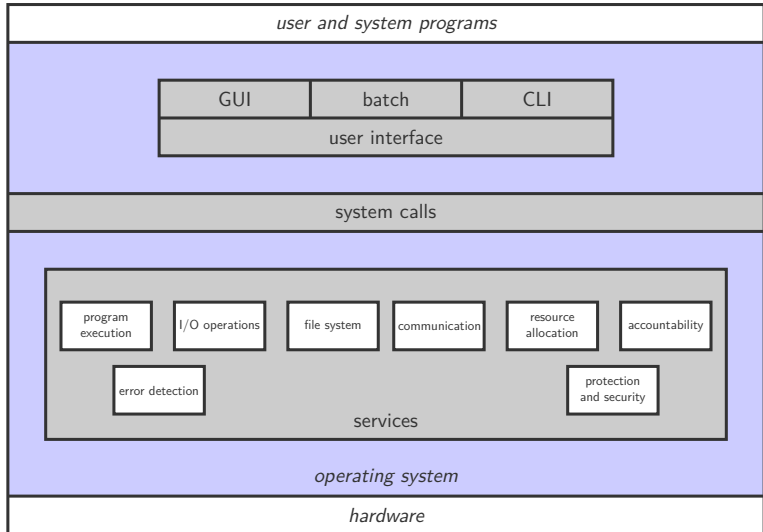
OS Services (1)

- The OS proposes two kinds of **services**
 - To the user
 - To the system itself
- Services offered **vary** by the OS

But common service classes are identifiable



OS Services (2)



System Call

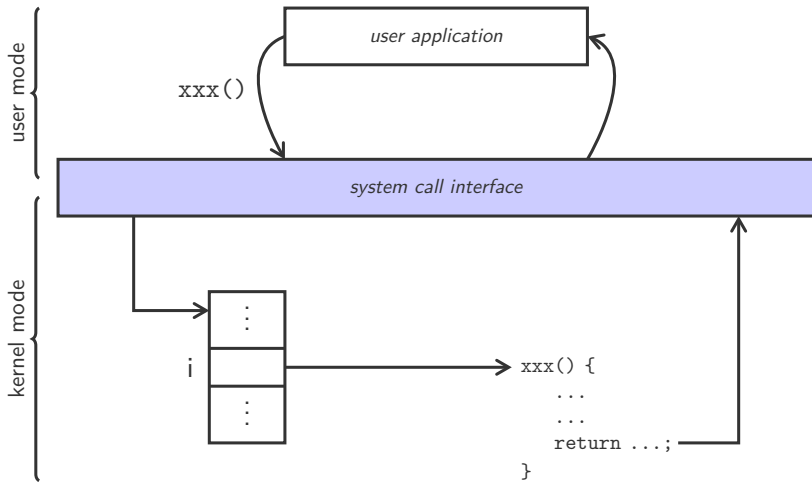
- Using the services of the OS through **system calls**
 - Interface to the services made available by the OS
 - Generally accessible as C/C++ routines
- Very **low level tasks** to write in assembly language

For example to make direct access to the hardware
- **Application Programming Interface (API)** (Win32, POSIX, Go...)
 - Encapsulate system calls sequence to avoid direct call
 - Access to an API from a library (`libc`, for example)

File Copy Example (1)

- 1 Get the name of the source file
 - Display a sentence on the standard output
 - Read the standard input
- 2 Get the name of the destination file
 - Display a sentence on the standard output
 - Read the standard input
- 3 Open the source file, create and open the destination file
- 4 As long as reading the file does not fail
 - Read a character from the source file
 - Write the character to the destination file
- 5 Close the source and destination files

System Call Execution (1)



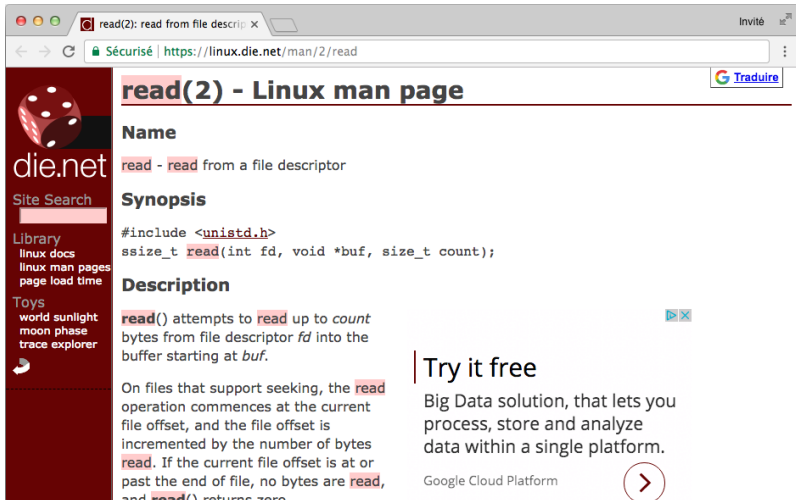
System Call Execution (2)

- Switching from user mode to **kernel mode**

The system call is executed in a privileged mode of the OS

- Management through a **system call interface**
 - Relay by intercepting system calls in the API
 - Use an associative table numbering the system calls
- The **user program** ignores system call details
 - It must respect the API and pass correct parameters
 - It must understand the returned value

System Call Documentation



The screenshot shows a web browser window with the address bar displaying `https://linux.die.net/man/2/read`. The page title is **read(2) - Linux man page**. On the left, there is a dark red sidebar with the **die.net** logo and a list of links: **Library** (linux docs, linux man pages, page load time), **Toys** (world sunlight, moon phase, trace explorer), and a Twitter icon. The main content area has a **Traduire** button in the top right. The page content includes:

- Name**: `read` - `read` from a file descriptor
- Synopsis**:

```
#include <unistd.h>
ssize_t read(int fd, void *buf, size_t count);
```
- Description**:

`read()` attempts to `read` up to `count` bytes from file descriptor `fd` into the buffer starting at `buf`.

On files that support seeking, the `read` operation commences at the current file offset, and the file offset is incremented by the number of bytes `read`. If the current file offset is at or past the end of file, no bytes are `read`, and `read()` returns zero.

On the right side of the description, there is a small icon with a blue 'x' and a red 'x'. Below the description, there is a section titled **Try it free** with the text: "Big Data solution, that lets you process, store and analyze data within a single platform." Below this text is the "Google Cloud Platform" logo and a red circular button with a white right-pointing arrow.

File Copy Example (2)

- Simplified version of a **file copy**

No error management at all

```
1 void copy(char *src, char*dst)
2 {
3     char buffer[BUFFSIZE];
4     int srcFile = open(src, O_RDONLY);
5     int dstFile = open(dst, O_WRONLY | O_CREAT, 0666);
6
7     if (srcFile != -1 && dstFile != -1) {
8         int r;
9         while ((r = read(srcFile, buffer, BUFFSIZE)))
10             write(dstFile, buffer, r);
11
12         close(srcFile);
13         close(dstFile);
14     }
15 }
```


System Call Categories

- Available **system calls** depend on the OS

Several system calls are common to most OS

- **Six main categories** of system calls can be identified
 - Process control
 - File management
 - Device management
 - IT maintenance
 - Communication
 - Protection

System Program

- Interface with the OS thanks to **system program**

System utility helps program development/execution

- Several **categories** of system programs
 - File management: `ls`, `mv`, `mkdir`, `rm`, etc.
 - Status information: `whoami`, `time`, `ps`, `top`, etc.
 - Text file editing: `vi`, `nano`, etc.
 - Support for programming languages: `cc`, `java`, etc.
 - Program loading and execution: `gdb`, `bash`, etc.
 - Communication: `ftp`, `mail`, `ssh`, etc.
 - Background services: `bg`, `screen`, `tmux`, etc.



Process and Memory

Process

- An OS can execute **several programs**
 - Resources sharing: CPU, memory, I/O devices, etc.
 - Compartmentalisation of programs
- A **process** is a running program
 - **Programme**: passive and on the disk
 - **Process**: active and in the memory



Process Abstraction

- A process is an **abstraction** for CPU resources

Executable software is organised in sequential processes

- Each process has its own **virtual processor**

It feels like it is the only one running on the processor

- No **timing** assumptions

A process can be stopped at any time by the OS

Process Management

- The OS must continuously **manage** the executed processes


On some systems, they are always active

- Several **operations** are made by the OS
 - Creation and deletion of processes
 - CPU allocation (*scheduling*) and other resource allocation
 - Synchronisation and communication
 - Deadlock management

Memory Structure

- A process is not only composed of the **code of the program**

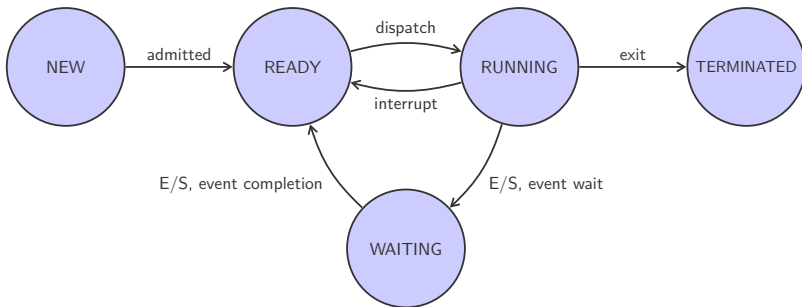
It is only one of the area used in memory



Five-State Model (1)

- A process can be in **different states** during its execution


Five main states are common to most OS



Five-State Model (2)

- The OS maintains several **data structures** to manage processes

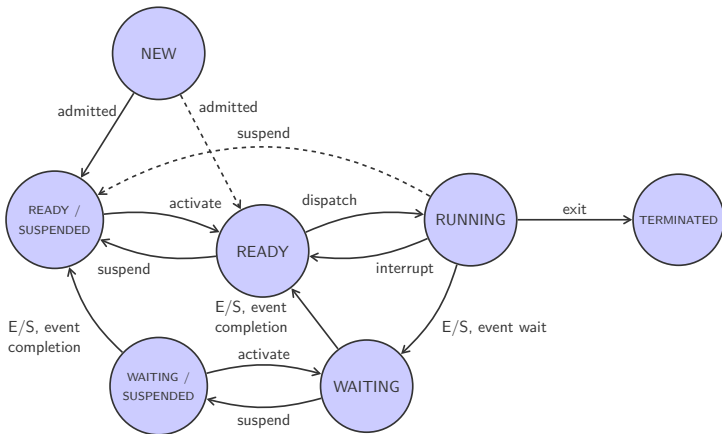
Several FIFO queues can be used to handle process execution



Process Suspension

- A process can be temporarily suspended

Moved from the memory to the disk (swapping)



Process Table

- The OS keeps a **process table** in memory

Each process has its own entry in this table

Process state
Process number
Ordinal counter
Registers
Memory limits
Open files list
...

Memory

- Program loaded in memory becomes a process

Disk → main memory

- Several algorithms are used to manage the memory

- 1 Machine level management

- 2 Pagation and segmentation

- The memory is just a large byte array

Each byte is uniquely identified by an address

Memory Unit

- When **executing an instruction** during a CPU cycle
 - **Load** instruction from the memory
 - **Store** instruction to the memory

- The **memory unit** sees a flow of memory addresses


The unit responds to the requests without any interpretation

- **Physical addresses** manipulation

Addresses of the bytes located in the main memory

Memory Protection (1)

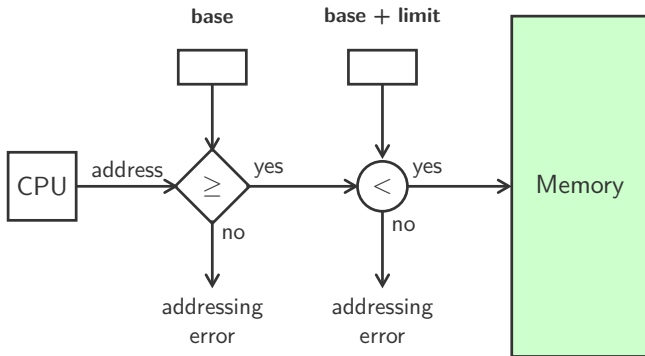
- Two registers are use to **delimit** the memory area
 - **base**: the smallest legal physical address
 - **limit**: the size of the memory area



Memory Protection (2)

- Only the OS can access the **base** and **limit** registers

Using a privileged instruction in kernel mode



Physical/Logical Address


- **Program and data** must be in memory, before
Size limitation corresponds to the physical memory
- Two separate and different **address spaces**
 - **Physical** addresses generated by the CPU
 - **Logical** addresses loaded in the memory-address register
This address is read by the memory unit

Memory Management Unit

- Hardware system to **convert** physical \rightarrow logical address

The user program only manipulates logical addresses


- A simple **MMU** just adds the value of a **relocation register**



Pagination

- Main memory divided into fixed size **frames**


Logical memory divided into pages with the same size



Logical
memory space

0	1
1	4
2	3
3	7

Page table



Physical
memory space

Address Computation (1)

- A **logical address** is a two-tuple

$\langle \text{page number}, \text{shift} \rangle$

- The OS holds one **page table** for each process

Entries with physical address of the page

- Two **registers** used to configure the page table


- Page-Table Base Register (PTBR)

- Page-Table Length Register (PTLR)

Address Computation (2)

- The address computation uses data from the **page table**

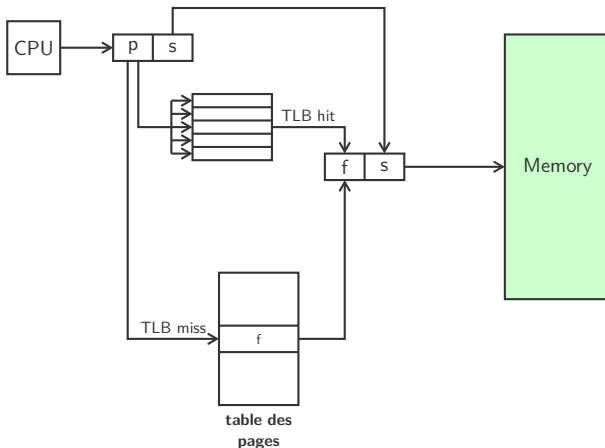
The shift is applied to the frame, once identified



Translation Look-Aside Buffer (TLB)

- The process can be speed up with a **fast cache memory**

Decreasing the number of main memory accesses



References

- Abraham Silberschatz, Peter B. Galvin, & Greg Gagne (2013). *Operating System Concepts*, John Wiley & Sons, ISBN: 978-1-11809-375-7.
- William Stallings (2017). *Operating Systems: Internals and Design Principles*, Pearson, ISBN: 978-1-29221-429-0.

Credits

- quapan, November 26, 2011, <https://www.flickr.com/photos/hinkelstone/7050697671>.
- Beao, September 9, 2009, http://commons.wikimedia.org/wiki/File:Personal_computer,_exploded_6.svg.
- Evan-Amos, February 19, 2017, <http://commons.wikimedia.org/wiki/File:PS4-Console-wDS4.jpg>.
- McZusatz, January 22, 2013, http://commons.wikimedia.org/wiki/File:Boeing_787-8_N787BA_cockpit.jpg.
- Jochembr, January 5, 2014, <http://commons.wikimedia.org/wiki/File:Snackomatic.jpg>.
- Dennis Skley, August 12, 2014, <https://www.flickr.com/photos/dskley/14711793077>.
- Jean-Pierre Dalbéra, September 23, 2014, <https://www.flickr.com/photos/dalbera/15766751411>.